

MacKenzie Scott’s giving, in QALYs

An evidence-weighted, auditable cost-effectiveness model of a \$26 billion portfolio

Max Ghenis*

MacKenzie Scott has disclosed \$26.39 billion in gifts to more than 2,500 organizations since 2019 — roughly a third of American megagiving in 2025 — yet almost none of it is denominated in outcomes. This paper prices the entire portfolio in quality-adjusted life-years. A Monte Carlo model allocates dollars across 14 intervention archetypes using Scott’s own gift-level database, imputes the undisclosed third from recipients’ pre-gift IRS 990 revenue, assigns each archetype a cost-per-QALY anchor from published causal estimates, and shrinks every effect toward zero in proportion to the causal credibility of its evidence. The best-guess estimate is a median of roughly 202,000 QALYs (90% interval 94,000–419,000), a blended \$150,000 per QALY, and a benefit–cost ratio of 4.9 at the U.S. government’s value per QALY. Three findings follow. First, impact is geographically concentrated: the 4.8% of dollars funding health and development delivered outside the United States produces 69% of estimated health impact. Second, gift size scales with recipient size to only the 0.41 power — a 10× larger organization receives about 2.5× more money. Third, the answer is dominated by evidence philosophy rather than arithmetic: restricting to randomized evidence collapses the median to ~9,000 QALYs while taking every study at face value raises it to ~415,000, a 45× range that dwarfs every other modeling choice. All parameters, code, data snapshots, and an AI-audit trail are public, and the model runs interactively in the browser.

Source: [Article Notebook](#)

1 Introduction

MacKenzie Scott gave away \$7 billion in 2025 — by the Indiana University Lilly Family School of Philanthropy’s accounting, about a third of every megagift in America that year (Indiana University Lilly Family School of Philanthropy 2026; Konish 2025). Her disclosed lifetime total is \$26.39 billion across 2,711 gifts to 2,545 organizations since 2019 (Yield Giving 2026). Almost none of it is denominated in health, and no published estimate prices the portfolio in outcomes of any kind.

*max@maxghenis.com. I am CEO of PolicyEngine and president of the Axiom Foundation; this paper is a personal research project and does not represent the views of either organization. Interactive version, code, data, and full audit trail: <https://maxghenis.com/mackenzie-scott-qaly> and <https://github.com/MaxGhenis/mackenzie-scott-qaly>.

This paper does so in the unit health economists use to compare a death averted against years lived in better health: the quality-adjusted life-year. The exercise is deliberately in the style of a charity evaluator’s cost-effectiveness analysis (GiveWell 2026), scaled from a program to a portfolio, with one addition the setting forces: because most of the portfolio funds interventions whose health effects have never been causally measured, the model makes its epistemics explicit. Every effect is shrunk toward zero in proportion to how well its supporting study identifies causation, the shrinkage schedule is a slider rather than a buried constant, and the headline is reported across the full range of evidence philosophies — from counting only randomized evidence to taking every published estimate at face value.

The paper makes four contributions. It provides, to my knowledge, the first portfolio-level cost-effectiveness estimate for a major philanthropist: a best-guess median of 202,000 QALYs from \$26.39 billion nominal (\$30.3B in 2026 dollars), a blended \$150,067 per QALY, and a benefit–cost ratio of 4.9 at the U.S. government’s value per QALY (U.S. Department of Health and Human Services, Office of the Assistant Secretary for Planning and Evaluation 2026). It documents extreme geographic concentration of modeled impact: the 4.8% of dollars funding global health and development delivered outside the United States accounts for 69.1% of estimated health impact. It reports a new empirical regularity in the gift data itself: across 1,313 disclosed gift–filing pairs, gift size scales with recipient pre-gift revenue to the power 0.41 (R^2 0.37) — Scott’s giving is far flatter across organization size than proportional. And it demonstrates a working method for auditable, AI-assisted cost-effectiveness modeling: the model was built with a frontier language model, reviewed and corrected in public across eleven revision rounds, audited by two independent model families against primary filings, and ships with versioned parameters, machine-checked citations, and a browser implementation that reruns the full Monte Carlo under any reader’s assumptions.

The unit of account deserves a caveat before any number. QALYs measure health. Most of Scott’s portfolio aims at other things — education, economic mobility, equity, civic infrastructure — whose primary benefits are not health benefits. A finding that a dollar of education philanthropy produces few QALYs is a statement about the lens, not a verdict on the gift. The estimates here are a lower bound on total wellbeing produced, measured along the one dimension with a mature comparative literature.

2 Data

2.1 The ledger

Scott’s essays disclose giving in dated tranches totaling \$26.39 billion nominal over 2020–2025. The model inflates each tranche to 2026 dollars with CPI-U, so dollars and cost-effectiveness evidence share one price level: \$30.3B total. All derivations below, and the replication repository, start from a verbatim snapshot of the Yield Giving gift database retrieved July 11, 2026 (Yield Giving 2026).

2.2 Allocation across causes

Scott publishes no aggregate dollars-by-cause breakdown, but she publishes better: a gift-level database listing every recipient organization, its self-reported focus areas, and — for 2,035 of 2,711 gifts (\$17.46B, about two-thirds of the nominal total) — the dollar amount. Each organization’s dollars are split equally across its focus areas; the database’s 53 areas map onto the model’s 14 intervention archetypes under a documented, sync-tested crosswalk.

The undisclosed third is imputed rather than dropped. Because the essays give each year’s total, the undisclosed residual per period is a known dollar amount; it is distributed across that period’s undisclosed gifts in proportion to each recipient’s pre-gift IRS Form 990 total revenue (ProPublica 2026), raised to an elasticity fit on the disclosed pairs. Pre-gift means the latest filing strictly before the gift year, since a Scott gift mechanically inflates receipt-year revenue. The fit is a finding in its own right: across 1,313 disclosed gift-filing pairs, log gift size rises with log revenue with slope 0.41 (R^2 0.37). A $10\times$ larger organization receives about $2.5\times$ more money, not $10\times$ more. Fuzzy name-to-EIN matches were audited by a third, smaller model against the live IRS API, with corrections committed as a reviewable overlay. The imputed third moves each cause share by at most about 1.5 percentage points — the disclosed two-thirds is representative.

2.3 Delivery geography

Each organization reports service locations. Weighting disclosed dollars by the share of an organization’s locations outside the United States, an estimated 20% of disclosed dollars (\$3.57 billion of \$17.46B) goes to organizations serving abroad, led by environment, health, and funding intermediaries. The model prices geography through a single channel: health-relevant dollars route to a global-health archetype in proportion to the recipient’s non-U.S. service share, because global health delivery is the one category with a published cost-per-QALY evidence base an order of magnitude below U.S. anchors. Education and climate dollars keep their U.S.-study anchors wherever they are delivered; if geography changes their effectiveness, the model does not capture it.

The coarsest geographic field is the bare label “global.” An audit of the 50 largest organizations reporting only that label — 87% of that bucket’s dollars — replaced the blanket assumption of fully-abroad delivery with fractions taken from each organization’s own filings and program reports, independently verified by a second model family, with per-organization sources committed to the repository. Many “global” organizations operate substantially in the United States: Planned Parenthood and Crisis Text Line, for example, list “global” service but are roughly 95% and 93% U.S. by their own reporting. The audited allocation routes 4.82% of the ledger to global health and development delivered abroad.

3 Model

Each Monte Carlo draw allocates the giving across archetypes (a Dirichlet centered on the derived shares), assigns each archetype a cost per QALY drawn from a distribution anchored in

published estimates, and multiplies the implied QALYs by two independent discounts. Causal credibility asks whether the measured effect is real. Each archetype carries the identification tier of its best supporting evidence, and the tier sets a Beta-distributed shrinkage factor toward the null: a lottery-randomized wealth shock (Cesarini et al. 2016) or the ACA Medicaid mortality quasi-experiment (Miller et al. 2021; Sommers 2017) keeps most of its effect; the community-health-center rollout design (Bailey and Goodman-Bacon 2015) and similar contested quasi-experiments keep less; associational and assumption-only categories are shrunk near zero. Realization asks whether a marginal philanthropic dollar delivers the studied effect at all — transport to different populations, overhead, and funging, net of the capacity effects of large unrestricted gifts that the Center for Effective Philanthropy documents for Scott’s recipients specifically (Center for Effective Philanthropy 2023). It is one global triangular draw with mode 0.80 on support 0.55–1.10.

The two discounts partition the doubt: transportability concerns live only in realization, identification concerns only in credibility, and cost-per-QALY priors are stated conditional on the effect being real and delivered. The tier ordering is the defensible part; the levels are judgment, proposed by a language model and reviewed by the author, and the results section reports the full consequence of disagreeing with them.

Health effects arrive in different natural units — lives saved, life-years, depression-free days — and are converted to discounted QALYs under a common set of conventions (about 65 remaining life-years at baseline utility 0.87, 3% annual discounting), with every parameter’s unit and dollar vintage machine-validated at load. The global-health frontier benchmark is handicapped like-for-like: GiveWell’s current program averages (GiveWell 2026), converted to roughly \$175–\$241 per discounted QALY-equivalent under this model’s own conventions, modeled as loguniform \$150–\$260, and subjected to the same realization and credibility discounts as Scott’s portfolio.

4 Results

Source: [Article Notebook](#)

Source: [Article Notebook](#)

At the best-guess defaults (100,000 draws), the portfolio produces a median of 201,923 QALYs, with a 90% interval of 94,000 to 419,000 and a mean of 222,060. The blended cost-effectiveness — dollars given over QALYs produced, at the median — is \$150,067 per QALY. Monetized at the U.S. Department of Health and Human Services’ value per QALY, the health alone repays the giving 4.9 times over (90% interval 2.1–11.0). Figure 1 shows the full distribution.

4.1 Where it comes from

Table 1 decomposes the mean estimate by cause. The result is the paper’s central fact: the global health and development slice — 4.8% of dollars — contributes 69.1% of estimated QALYs, at a realized cost of about \$10,000 per QALY against \$350,000–\$450,000 for the large U.S. buckets. The 1,200-fold spread in realized cost-effectiveness within a single donor’s portfolio restates, in

Mackenzie Scott's lifetime giving — estimated health impact (QALYs)



Figure 1: Distribution of total QALYs across 100,000 Monte Carlo draws at the best-guess defaults. The dashed line marks the mean; the solid line the median.

one column, the central empirical claim of the effective-altruism literature: the cheapest health in the world is not bought where most philanthropic dollars are spent.

Source: [Article Notebook](#)

Table 1: Dollars and estimated QALYs by cause, best-guess defaults. Dollars apply the allocation centers to the 2026-dollar total; QALYs are mean contributions (only means are additive). \$/QALY is the realized ratio after credibility and realization discounts, not the raw study anchor.

Cause	Dollars	% of \$	QALYs	% of QALYs	\$/QALY	Evidence
Global health & development (non-US delivery)	\$1.45B	5%	153,469	69%	\$9,477	strong_quasi
Education (HBCUs, comm. college, scholarships)	\$5.70B	19%	18,939	9%	\$300,788	moderate_quasi
Health - insurance & access	\$1.12B	4%	9,977	4%	\$112,381	strong_quasi
Health - mental & behavioral	\$667M	2%	8,810	4%	\$75,673	randomized
Economic security - workforce & mobility	\$4.06B	13%	6,632	3%	\$612,212	observational
Economic security - housing & homelessness	\$1.09B	4%	5,918	3%	\$184,341	moderate_quasi
Health - community health centers	\$1.12B	4%	4,062	2%	\$276,005	moderate_quasi
Other / general community	\$1.67B	6%	3,708	2%	\$449,423	observational
Environment / climate	\$2.48B	8%	3,193	1%	\$778,250	projection

Table 1: Dollars and estimated QALYs by cause, best-guess defaults. Dollars apply the allocation centers to the 2026-dollar total; QALYs are mean contributions (only means are additive). \$/QALY is the realized ratio after credibility and realization discounts, not the raw study anchor.

Cause	Dollars	% of \$	QALYs	% of QALYs	\$/QALY	Evidence
Economic security - cash & financial	\$879M	3%	2,880	1%	\$305,164	strong_quasi
Economic security - food security	\$576M	2%	2,093	1%	\$275,096	observa-
Equity & justice	\$5.21B	17%	1,763	1%	\$3M	tional
Civic / democracy / philanthropy infra	\$3.24B	11%	531	<1%	\$6M	assump-
Arts & culture	\$1.03B	3%	85	<1%	\$12M	tion
Total	\$30.30B	100%	222,060	100%	\$136,458	assump-

Source: [Article Notebook](#)

Source: [Article Notebook](#)

Table 2 decomposes the same run by delivery geography. The table is a by-construction decomposition, not evidence of regional differences: within each cause the model prices every delivery location identically — the health routing to global anchors is its only geographic pricing — so a cause’s QALYs follow its dollars across regions. Under that construction, Sub-Saharan Africa receives 6% of dollars and 30% of estimated QALYs at \$27,031 per QALY; nationally-scoped U.S. giving receives 27% of dollars and 9% of QALYs at \$403,000.

Source: [Article Notebook](#)

Table 2: The same run decomposed by delivery geography. Asterisked rows are reporting granularity (national or multi-region listings), not places. Geography derives from each organization’s reported service locations, equal-split, with audited fractions for the 50 largest “global”-listed organizations.

Region	Dollars	% of \$	QALYs	% of QALYs	\$/QALY
Sub-Saharan Africa	\$1.81B	6%	67,007	30%	\$27,031
South Asia	\$1.01B	3%	32,006	14%	\$31,577
Global (multi-region) *	\$1.53B	5%	28,288	13%	\$54,164
US — national / unspecified *	\$8.24B	27%	20,443	9%	\$403,000
Latin America & Caribbean	\$1.11B	4%	17,472	8%	\$63,620
US — South	\$5.96B	20%	16,041	7%	\$371,678
US — West	\$4.02B	13%	10,966	5%	\$366,611
US — Midwest	\$2.75B	9%	7,828	4%	\$351,406
US — Northeast	\$2.20B	7%	5,631	3%	\$390,341

Table 2: The same run decomposed by delivery geography. Asterisked rows are reporting granularity (national or multi-region listings), not places. Geography derives from each organization’s reported service locations, equal-split, with audited fractions for the 50 largest “global”-listed organizations.

Region	Dollars	% of \$	QALYs	% of QALYs	\$/QALY
Europe & Central Asia	\$299M	1%	5,453	2%	\$54,776
Middle East & North Africa	\$153M	1%	4,177	2%	\$36,670
East Asia & Pacific	\$335M	1%	3,339	2%	\$100,390
North America (unspecified) *	\$698M	2%	1,535	1%	\$454,724
Other / unmapped *	\$40M	<1%	1,435	1%	\$27,621
Canada	\$142M	<1%	439	<1%	\$323,208

Source: [Article Notebook](#)

4.2 Against the global-health frontier

Handicapped with the same realization and credibility discounts, the GiveWell-derived frontier benchmark delivers roughly $521\times$ more health per marginal dollar than the blended portfolio; the same sum directed entirely to the frontier would imply on the order of 105.5 million QALYs. That figure is a marginal comparison, not an attainable counterfactual: frontier-priced opportunities are scarce, GiveWell directs hundreds of millions — not tens of billions — per year at these prices, and funding beyond identified room drives toward cash-transfer cost-effectiveness. Redeploying the full portfolio would compress the multiple toward a floor plausibly under $20\text{--}40\times$ (GiveWell 2024). The honest statement is directional: the frontier multiple is several hundred at today’s margin and remains large under any plausible account of diminishing returns.

5 Robustness

5.1 Evidence philosophy dominates everything else

Source: [Article Notebook](#)




Figure 2: Median total QALYs across the evidence-stance range, 100,000 draws per point. The stance interpolates every tier’s credibility mean from an RCT-only floor (non-randomized tiers shrunk to 0.01) through the drafted best-guess tiers (0.5) to face value (1.0).

Source: [Article Notebook](#)

The dominant sensitivity is not a parameter but a philosophy. Restricting credit to randomized evidence collapses the median to 9,114 QALYs — dominated by the one lottery-randomized bucket — while taking every cited effect at face value raises it to 414,264 (Figure 2). That 45-fold range dwarfs the spread from any distributional parameter, and the gap between the face-value and best-guess estimates is a direct reading of how much of claimed philanthropic impact rests on evidence that would not survive a causal-inference seminar. The Spearman tornado (Figure 3) tells the same story within the default stance: the top three drivers are the global-health allocation share, the realization factor, and global-health credibility.

5.2 The audit moved the headline by less than 2%

The geography audit is the largest data correction the model has absorbed since publication, and its effect is small: recentering the allocation on audited service geographies moved the abroad share from 4.93% to 4.82%, the median from about 205,000 to 201,923, and the blended figure from about \$148,000 to \$150,067 per QALY. Both allocations ship with the model, the browser version exposes them as a version toggle, and the pre-audit parameters remain pinned at tag v2026.07.20 (audited: v2026.07.22).




Figure 3: What drives the spread: Spearman rank correlations between each input draw and total QALYs at the best-guess defaults.

5.3 Corrections history

The model reached this state through eleven public revision rounds, and the trajectory is itself evidence about the method. Cross-model review rounds cut the headline from roughly 98,000 to 87,000 QALYs (dollar-vintage normalization, a life-year/QALY conversion made explicit, a re-derived frontier benchmark, one re-attributed citation). Replacing an assumed allocation with the derivation from Scott’s own database cut it to roughly 70,000. A reader’s correction — health dollars had been priced at U.S. anchors regardless of delivery location — tripled it to roughly 205,000, and the subsequent two-model geography audit settled it at 202,000. Every correction is documented in the repository and summarized on the public tool; none was caught by the author unaided, and all but one were caught by models or readers rather than the author.

6 Construction and auditability

The model was built with a frontier language model (Anthropic’s Claude), with the author reviewing every parameter and assumption; adversarial review by a second model family (OpenAI’s), plus reader corrections, produced the revision history above. Three mechanisms make the result checkable rather than merely stated. Parameters live in one typed file in which every value carries its unit, dollar vintage, and inline citation, machine-validated at load; an annotated bibliography maps each quantitative input to its primary source and price-level conversion. Sync tests force the published web parameters, allocation shares, geography exports, and this manuscript’s inputs to equal fresh derivations from the committed data snapshots, so prose and model cannot drift apart. And the audit trail is data: the 990 match audit and the 50-organization geography audit are committed as per-row overlay files with sources and confidence grades, applied by the pipeline rather than folded silently into parameters.

The browser implementation is a checked TypeScript port that reruns the full Monte Carlo client-side; every figure in this paper can be reproduced, and every assumption overridden, at maxghenis.com/mackenzie-scott-qaly. The repository, including this manuscript’s source, is at github.com/MaxGhenis/mackenzie-scott-qaly.

7 Limitations

The credibility tiers and realization distribution are judgment priors with no literature calibration; the ordering is defensible, the levels are not measured, and the stance sweep is the honest bound on their influence. The QALY lens counts only health, so the estimate is a floor on total wellbeing, and the causes where Scott concentrates — education, mobility, equity — are precisely those whose main benefits the lens excludes. Credibility shrinks toward zero but never below it, so the model places no mass on harmful giving. Allocation splits each organization’s dollars equally across its focus areas and service locations; both are coarse, and the geography audit covers only the 50 largest “global”-listed organizations. Realization is one global draw, so it cannot vary by cause. Archetype anchors treat heterogeneous organizations as draws from one distribution. Within each cause, delivery regions are priced identically, so the geographic decomposition exposes assumptions rather than measuring regional differences. The undisclosed third of dollars

is imputed from a fitted elasticity, not observed. Finally, several anchors rest on a small number of quasi-experimental papers whose external validity to Scott’s grantees is exactly what the realization factor is guessing at.

8 Conclusion

Measured in health, MacKenzie Scott’s \$26 billion buys a median of about 202,000 quality-adjusted life-years — several times its cost at the government’s own valuation — and most of that health comes from the twentieth of her dollars delivered in low-income countries. The number is a model output under stated priors, not a fact, and the model’s most important output is not the point estimate but the machinery around it: explicit evidence weights, a 45-fold stance range that any reader can traverse, versioned corrections, and audits committed as data. Portfolio-scale philanthropy is measurable in outcomes without waiting for donors to publish outcome data — their own gift lists, recipients’ tax filings, and the published causal literature suffice — and the same machinery applies to any large giver. The invitation of the interactive version is the paper’s real conclusion: disagree with a parameter, move it, and watch what survives.

Acknowledgments

I thank readers on LinkedIn and the EA Forum whose corrections materially improved the model, including the reader who identified the original mispricing of internationally delivered health dollars. The model, audits, and this manuscript were built with Anthropic’s Claude, with additional adversarial review by OpenAI models; all judgments and errors are mine. This is a personal research project and does not represent the views of PolicyEngine or the Axiom Foundation.

References

Source: [Article Notebook](#)

- Bailey, Martha J., and Andrew Goodman-Bacon. 2015. “The War on Poverty’s Experiment in Public Medicine: Community Health Centers and the Mortality of Older Americans.” *American Economic Review* 105 (3): 1067–104. <https://doi.org/10.1257/aer.20120070>.
- Center for Effective Philanthropy. 2023. Emerging Impacts: The Effects of MacKenzie Scott’s Large, Unrestricted Gifts (Year 2). Center for Effective Philanthropy. https://cep.org/wp-content/uploads/2023/11/BigGiftsStudy_Report_Y2_FNL.pdf.
- Cesarini, David, Erik Lindqvist, Robert Östling, and Björn Wallace. 2016. “Wealth, Health, and Child Development: Evidence from Administrative Data on Swedish Lottery Players.” *Quarterly Journal of Economics* 131 (2): 687–738. <https://doi.org/10.1093/qje/qjw001>.

- GiveWell. 2024. Re-Evaluating the Impact of Unconditional Cash Transfers. <https://blog.givewell.org/2024/11/12/evaluating-the-impact-of-unconditional-cash-transfers/>.
- GiveWell. 2026. Impact Estimates. <https://www.givewell.org/impact-estimates>.
- Indiana University Lilly Family School of Philanthropy. 2026. Giving USA 2026: The Annual Report on Philanthropy for the Year 2025. Giving USA Foundation.
- Konish, Lorie. 2025. MacKenzie Scott Revealed Her Total Charitable Donations for 2025. <https://www.cnbc.com/2025/12/13/mackenzie-scott-revealed-her-total-charitable-donations-for-2025.html>.
- Miller, Sarah, Norman Johnson, and Laura R. Wherry. 2021. “Medicaid and Mortality: New Evidence from Linked Survey and Administrative Data.” *Quarterly Journal of Economics* 136 (3): 1783–829. <https://doi.org/10.1093/qje/qjab004>.
- ProPublica. 2026. Nonprofit Explorer. <https://projects.propublica.org/nonprofits/>.
- Sommers, Benjamin D. 2017. “State Medicaid Expansions and Mortality, Revisited: A Cost-Benefit Analysis.” *American Journal of Health Economics* 3 (3): 392–421. https://doi.org/10.1162/ajhe_a_00080.
- U.S. Department of Health and Human Services, Office of the Assistant Secretary for Planning and Evaluation. 2026. Standard Values for Regulatory Impact Analysis. ASPE. <https://aspe.hhs.gov/reports/standard-ria-values>.
- Yield Giving. 2026. Gifts Database. <https://yieldgiving.com/gifts>.